

# A Hybrid Framework for Fluid Flow Simulations: Combining SPH with Machine Learning

**Rene Winchenbach**, Nils Thuerey


Technical University Munich

SPHERIC 2023

Rhodes, 28. June 2023



<https://tinyurl.com/sphericDemo>



# Why Machine Learning?

We distinguish two motivational approaches:

1. Computational Performance
2. Methodological limitations

The former is always helpful as faster simulations are never bad


The latter is much more useful in practice

Design optimization common in CFD


Goal: Minimize cost function w.r.t. shape

SPH is not easily differentiable


ML surrogates are automatically differentiable




(a) Direction (GD-M)  
 $r = 1.19$  [1.184, 1.201]



(b) 2 Pools (GD-M)  
 $r = 0.90$  [0.897, 0.909]



(c) 3 Pools (GD-M)  
 $r = 0.87$  [0.844, 0.887]



# Graph Convolutions and SPH

SPH is a graph convolution

A particle is a vertex of an implicit graph

Neighborhoods form the connectivity of the graph

SPH interpolations are message passing steps:


$$\langle A(x_i) \rangle = \sum_j \frac{m_j}{\rho_j} A_j W(x_j - x_i)$$

A (continuous) graph convolution of feature  $f$  with  $g$ :

$$(f \star g)(x_i) = \sum_j f_j g(x_j - x_i)$$

Where  $f_j = \frac{m_j}{\rho_j} A_j$  and  $g = W$  leads to an SPH interpolation step

GNNs *simply* learn  $g$



# What **we** will do now

The simplest SPH interpolation is a density summation step:

$$\rho_i = \sum_j m_j W_{ij}$$

We are going to learn this in 1D and we are going to do it **live**

For this case we know what the ideal graph convolution looks like

Ideally  $g$  should be equal to  $W$  (for our case a Wendland 2 kernel function)

# PyTorch

PyTorch is a library for python developed by Meta

It is focused on tensors, which are ubiquitous in neural networks

Interface very similar to numpy for mathematical operations

Builtin neural network layer modules

Ecosystem of libraries, e.g., for GNNs

Supports automatic differentiation

Supports automatic parallelization

```

1 import torch
2
3 # Create two tensors with requires_grad=True
4 x = torch.tensor(2.0, requires_grad=True)
5 y = torch.tensor(3.0, requires_grad=True)
6
7 # Perform some operations on the tensors  $z = x^2 + y^2$ 
8 z = x * x + y * y
9
10 # Call backward() to compute gradients
11 z.backward()
12
13 # Check the gradients
14 print(f"Gradient of z with respect to x: {x.grad}")
15 print(f"Gradient of z with respect to y: {y.grad}")

```

$$\frac{\partial}{\partial x} z \Big|_{x=2, y=3}$$

$$\frac{\partial}{\partial y} z \Big|_{x=2, y=3}$$

# Google Colab

An important aspect of research and education is accessibility

Not everyone has powerful hardware to run simulations

There are many cloud based platforms to run computations on

Google Colab allows for code to be ran on GPUs by anyone with a Google account for free

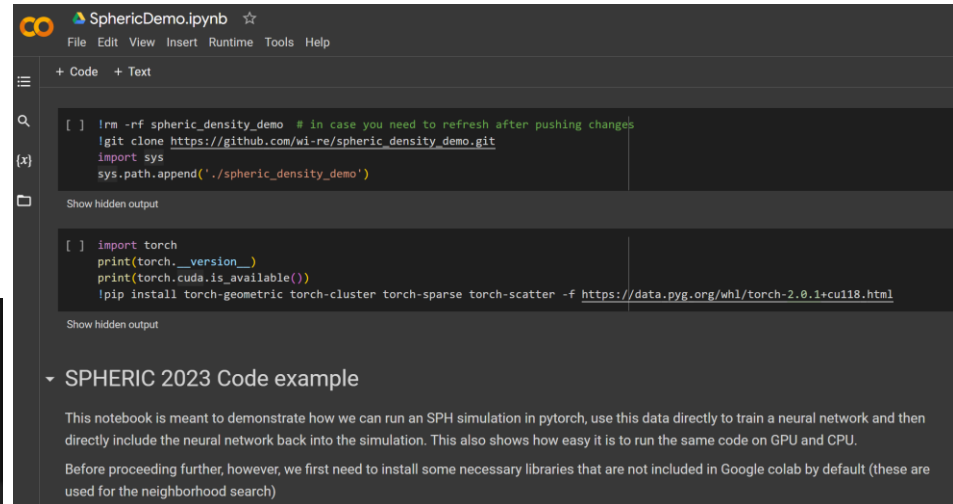
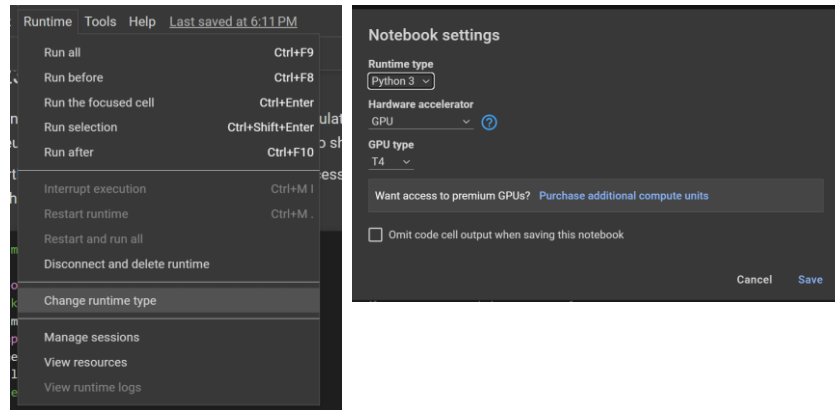
Interface very similar to Jupyter Notebooks

Colab notebooks can be shared via simple links and require no local installations

But there are limitations (certainly to the free version)

# Google Colab

This is what the demo should look like  
First we make sure that we have GPU support



Note that the notebook can also be run without any GPU acceleration and instead using CPU parallelization

# Demo I – Installing Prerequisites

To get going we need to install some python modules and libraries:

We can directly clone a Github repository with our own modules:

```
[ ] !rm -rf spheric_density_demo # in case you need to refresh after pushing changes
!git clone https://github.com/wi-re/spheric_density_demo.git
import sys
sys.path.append('./spheric_density_demo')

Cloning into 'spheric_density_demo'...
remote: Enumerating objects: 29, done.
remote: Counting objects: 100% (29/29), done.
remote: Compressing objects: 100% (23/23), done.
remote: Total 29 (delta 11), reused 20 (delta 5), pack-reused 0
Unpacking objects: 100% (29/29), 4.80 MiB | 8.48 MiB/s, done.
```

And then install PyTorch Geometric for GNN features using pip (make sure the versions match)

```
[ ] import torch
print(torch.__version__)
print(torch.cuda.is_available())
!pip install torch-geometric torch-cluster torch-sparse torch-scatter -f https://data.pyg.org/whl/torch-2.0.1+cu118.html
```



## Demo II – Loading Modules

Next we import the python modules we need for this demo and are ready to go

```
[ ] # sph related imports
    from sph import *
    from perlin import *
    # neural network related imports
    from torch.optim import Adam
    from rbfConv import *
    from torch_geometric.loader import DataLoader
    from trainingHelper import *
    # plotting/UI related imports

    from plotting import *
    import matplotlib as mpl
    plt.style.use('dark_background')
    cmap = mpl.colormaps['viridis']
    from tqdm.notebook import trange, tqdm
    from IPython.display import display, Latex
```

# Governing Equations

Compressible 1D SPH simulation  $\frac{d\mathbf{u}}{dt} = \frac{\nabla p}{\rho} + \left(\frac{d\mathbf{u}}{dt}\right)_{diss}$

Pressure based on ideal gas EOS  $p_i = \kappa (\rho_i - \rho_0)$ ;  $\frac{\nabla p_i}{\rho_i} = \sum_j m_j \left( \frac{p_i}{\rho_i^2} + \frac{p_j}{\rho_j^2} \right) \nabla_i W_{ij}$


Kinematic viscosity based on standard SPH formulations  $\left(\frac{d\mathbf{u}_i}{dt}\right)_{diss} = -\sum_j m_j \frac{-\alpha c_s \mu_{ij} + \beta \mu_{ij}^2}{(\rho_i + \rho_j)/2} \nabla_i W_{ij}$ ;  $\mu_{ij} = h \frac{\mathbf{u}_{ij} \cdot \mathbf{r}_{ij}}{r_{ij} + \epsilon h^2}$

Density summation based  $\rho_i = \sum_j m_j W_{ij}$

RK4 time integration

Explicit fixed time stepping

Periodic Boundary Conditions



# Simulation Initialization

The simulation is initialized based on a desired density profile defined using a PDF


The default setup is a sinusoidal density profile

This profile creates a very regular simulation outcome that is easy to understand

The demo also comes with many other initial PDFs, including random ones


Simulation parameters:

```
# simulation parameters
minDomain = -1 # minimum domain, leave at -1 for the most part
maxDomain = 1 # maximum domain, leave at 1 for the most part
# change base area to change initial starting density
baseArea = 2 / numParticles * 2
particleRadius = baseArea / 2.0
# change particle support to make simulation more/less smooth
particleSupport = particleRadius * 8.
# SPH parameters
xsphConstant = 0.0
diffusionAlpha = 1. # kinematic viscosity coefficient
diffusionBeta = 2.
kappa = 10 # EOS kappa term
restDensity = 1000 # EOS rest density term
dt = 1e-3 # fixed global timestep
c0 = 10 # speed of sound used in kinematic viscosity
```




# Example Simulations


## Sinusoidal




## Sawtooth



## Rectangular



## Random



# Demo III – Running the Simulation


Next we run the simulation with 4 substeps

We also write out a lot of data to process it later

```
[ ] # run the simulation using RK4
for i in tqdm(range(timesteps)):
    # Compute state for substep 1
    v1 = torch.clone(fluidVelocities)
    # RK4 substep 1
    dudt_k1, dxdt_k1, fluidDensity, fluidPressure = computeUpdate(fluidPositions, fluidVelocities, fluidAreas, minDomain, maxDomain, kappa, restDensity, diffusionAlpha, diffusionBeta, c0, xsphConstant, particleSupport, dt)
    # Compute state for substep 2
    x_k1 = fluidPositions + 0.5 * dt * dxdt_k1
    x_k1[x_k1 < minDomain] += maxDomain - minDomain
    x_k1[x_k1 > maxDomain] -= maxDomain - minDomain
    u_k1 = fluidVelocities + 0.5 * dt * dudt_k1
    # RK4 substep 2
    dudt_k2, dxdt_k2, _, _ = computeUpdate(x_k1, u_k1, fluidAreas, minDomain, maxDomain, kappa, restDensity, diffusionAlpha, diffusionBeta, c0, xsphConstant, particleSupport, 0.5 * dt)
    # Compute state for substep 2
    x_k2 = fluidPositions + 0.5 * dt * dxdt_k2
    x_k2[x_k2 < minDomain] += maxDomain - minDomain
    x_k2[x_k2 > maxDomain] -= maxDomain - minDomain
    u_k2 = fluidVelocities + 0.5 * dt * dudt_k2
    # RK4 substep 3
    dudt_k3, dxdt_k3, _, _ = computeUpdate(x_k2, u_k2, fluidAreas, minDomain, maxDomain, kappa, restDensity, diffusionAlpha, diffusionBeta, c0, xsphConstant, particleSupport, 0.5 * dt)
    # Compute state for substep 4
    x_k3 = fluidPositions + dt * dxdt_k3
    x_k3[x_k3 < minDomain] += maxDomain - minDomain
    x_k3[x_k3 > maxDomain] -= maxDomain - minDomain
    u_k3 = fluidVelocities + dt * dudt_k3
    # RK4 substep 4
    dudt_k4, dxdt_k4, _, _ = computeUpdate(x_k3, u_k3, fluidAreas, minDomain, maxDomain, kappa, restDensity, diffusionAlpha, diffusionBeta, c0, xsphConstant, particleSupport, dt)
    # RK substeps done, store current simulation state for later processing/learning. density and pressure are based on substep 1 (i.e., the starting point for this timestep)
    simulationStates.append(torch.stack([fluidPositions, fluidVelocities, fluidDensity, fluidPressure, dt/6 * (dudt_k1 + 2* dudt_k2 + 2* dudt_k3 + dudt_k4), dudt_k1, dudt_k2, dudt_k3, dudt_k4, fluidAreas]))
    # time integration using RK4 for velocity
    fluidVelocities = fluidVelocities + dt * dudt_k1 # semi implicit euler mode
    fluidVelocities = fluidVelocities + dt/6 * (dudt_k1 + 2* dudt_k2 + 2* dudt_k3 + dudt_k4)
    fluidPositions = fluidPositions + dt * fluidVelocities
    # enforce periodic boundary conditions
    fluidPositions[fluidPositions < minDomain] += maxDomain - minDomain
    fluidPositions[fluidPositions > maxDomain] -= maxDomain - minDomain
    # After the simulation has run we stack all the states into one large array for easier slicing and analysis
    simulationStates = torch.stack(simulationStates)
```

# Demo IV – Simulation Results

The results should look like this:



# Datasets and You

An important aspect of neural networks is the data generation/acquisition

For this demo we only have one simulation setup for training and testing

This is not ideal but suffices for this demonstration

Ideally the dataset should be as varied as possible

This includes using data augmentation to include simulation states that aren't possible in SPH

Batching the training can improve training stability but increases memory pressure

```
[ ] ignoredTimesteps = 256
    batchSize = 4 # training batch size
    # Training done on all timesteps except the last ignoredTimesteps
    timestamps = np.arange(0, simulationStates.shape[0] - ignoredTimesteps)
    testBatch = np.arange(len(simulationStates) - ignoredTimesteps, len(simulationStates))


    # create pytorch dataloader (via pytorch geometric for convenience)
    dataLoader = DataLoader(timestamps, shuffle=True, batch_size = batchSize).batch_sampler
    dataIter = iter(dataLoader)
```

## Demo V – Neural Network Setup

Next we setup a single layer neural network with empirically chosen hyperparameters:

```
[ ] # Hyperparameters for the NN
    lr = 1e-1 # Learning rate
    iterations = 1000 # update iterations per epoch
    epochs = 5 # total number of epochs, LR is halved every epoch
    n = 15 # number of weights per continuous convolution
    basis = 'linear' # basis for the convolution, set to linear for CConv
    computeBatchSize = 128 # higher number = faster processing but more memory consumption (not relevant for small simulations)
    windowFn = None # getWindowFunction('Wendland2_1D')
    normalized = False # rbf normalization
```

We can then look at the initialized neural network and what the convolution looks like:



Which is certainly not a compact gaussian.



# Demo VI – Neural Network Training

We now train the neural network:

```
[ ] # create progress bar and arrays to hold the learning progress
pb = tqdm(range(epochs * iterations))
lossArray = []
batches = []
weights = []
testing = []
# the actual learning process
for epoch in range(epochs):
    losses = []
    b, l, w = processDataLoaderIter(pb, iterations, epoch, lr, \
                                   dataLoader, dataIter, batchSize, model, optimizer, \
                                   simulationStates, minDomain, maxDomain, particleSupport, \
                                   lossFunction, getFeatures, getGroundTruth, None, \
                                   train = True, prefix = '', augmentAngle = False, augmentJitter = False, jitterAmount = 0.01)

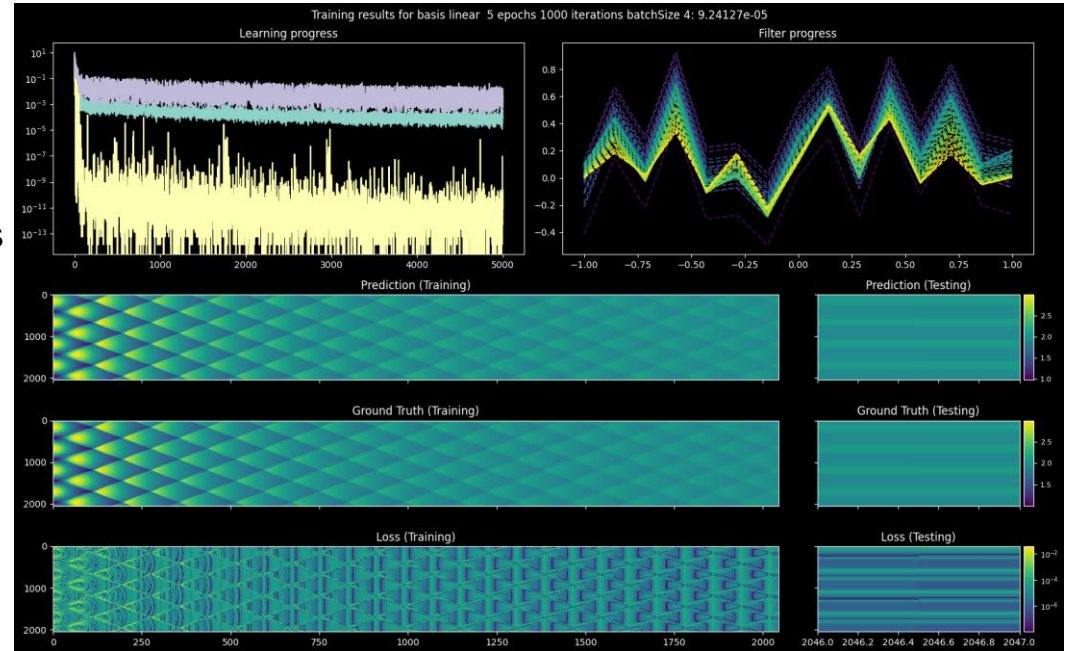
    lossArray.append(l)
    batches.append(b)
    weights.append(w)
    if True: #epoch % 1 == 0 and epoch > 0:
        lr = lr * 0.5
        for param_group in optimizer.param_groups:
            param_group['lr'] = 0.5 * param_group['lr']
```

The learning rate is halved every 1000 weight updates to improve training stability

We also write out a lot of information to visualize the process

# Demo VII – Neural Network Evaluation

The training results should look like this:  
 Prediction and ground truth are very similar  
 The learned convolution is still not gaussian  
 The training loss is small compared to the numerical precision we get from 32 bit floats  
 So, is this good enough?



# Building a hybrid simulation

As the trained network is a graph convolution, including this into the simulation is straight forward:



```
# SPH simulation step, returns dudt, dxdx as well as current density and pressure
def computeUpdateML(model, fluidPositions, fluidVelocities, fluidAreas, minDomain, maxDomain, kappa, restDensity, diffusionAlpha, diffusionBeta, c0, xsphCoefficient, particleSupport, dt):
    # 1. Create ghost particles for our boundary conditions
    ghostPositions = createGhostParticles(fluidPositions, minDomain, maxDomain)
    # 2. Find neighborhoods of all particles:
    fluidNeighbors, fluidRadialDistances, fluidDistances = findNeighborhoods(fluidPositions, ghostPositions, particleSupport)

    # ----- #
    # 3. Compute \rho using the trained neural network
    features = getFeatures(fluidPositions, fluidAreas, fluidVelocities, None)
    with torch.no_grad():
        fluidDensity = model((features[:,None], features[:,None]), fluidNeighbors, fluidRadialDistances[:,None] * torch.sign(fluidDistances[:,None]))[:,0]
    # ----- #

    # 4. Compute the pressure of each particle using an ideal gas EOS
    fluidPressure = (fluidDensity - 1.0) * kappa * restDensity
    # 5. Compute the XSPH term and apply it to the particle velocities:
    xsphUpdate = computeXSPH(fluidPositions, fluidVelocities, fluidDensity, fluidAreas, particleSupport, xsphCoefficient, fluidNeighbors, fluidRadialDistances)
    # 6. Compute pressure forces and resulting acceleration
    fluidPressureForces = computePressureForces(fluidPositions, fluidDensity, fluidPressure, fluidAreas, particleSupport, restDensity, fluidNeighbors, fluidRadialDistances, fluidDistances)
    fluidAccel = fluidPressureForces
    # 7. Compute kinematic viscosity
    laminarViscosity = computeDiffusion(fluidPositions, fluidVelocities, fluidAreas, fluidDensity, particleSupport, restDensity, diffusionAlpha, diffusionBeta, c0, fluidNeighbors, fluidRadialDistances, fluidDistances)
    fluidAccel += xsphUpdate / dt + laminarViscosity
    return fluidAccel, fluidVelocities, fluidDensity, fluidPressure
```

# Demo VII – Hybrid Simulation Results

After doing this we get the hybrid simulation on the right compared to the ground truth on the left



# Conclusions

Machine Learning is probably inevitable even in CFD  
There are many opportunities for research in this direction  
Making research more accessible is important  
Trusting neural networks is a major challenge in many fields  
Conducting open research that can be verified and repeated  
even more important than in classical research



<https://tinyurl.com/sphericDemo>

Scan the QR Code on the right to go to the Google Colab notebook and try it out yourself!  
Or go to [https://github.com/wi-re/spheric\\_density\\_demo.git](https://github.com/wi-re/spheric_density_demo.git) also has the presentation slides